


TensorFold · Swift 1.5 solo comparison


M5 Max Mac Studio · 1 request · previous vs updated configuration

Previous Swift 1.5 configuration Updated Swift 1.5 configuration

Generation speed · higher is better



Effective prompt throughput · higher is better



Updated vs previous

Generation change

+4.3%

+3.3%

+2.5%

-10.1%

+0.2%

+17.0%

Prompt throughput change

+2.9%

+2.3%

+1.9%

+2.8%

+2.6%

+1.2%

Previous trials: 3 at 20K/40K/60K; 1 at 80K/100K/128K. Updated trials: 3 at 20K-100K; 1 at 128K.

Points: medians; whiskers: observed min-max, not confidence intervals. Solo outputs: 1,024 tokens; zero cache hits.

The blue series matches the published Swift chart; its 128K point uses the prior higher-memory rerun.

Historical configuration comparison, not a controlled A/B. Model revision, sampling and output limit match.

Effective prompt throughput = uncached input tokens ÷ (queue + prefill time), not isolated GPU prefill.