


TensorFold · Swift 1.5 4-bit

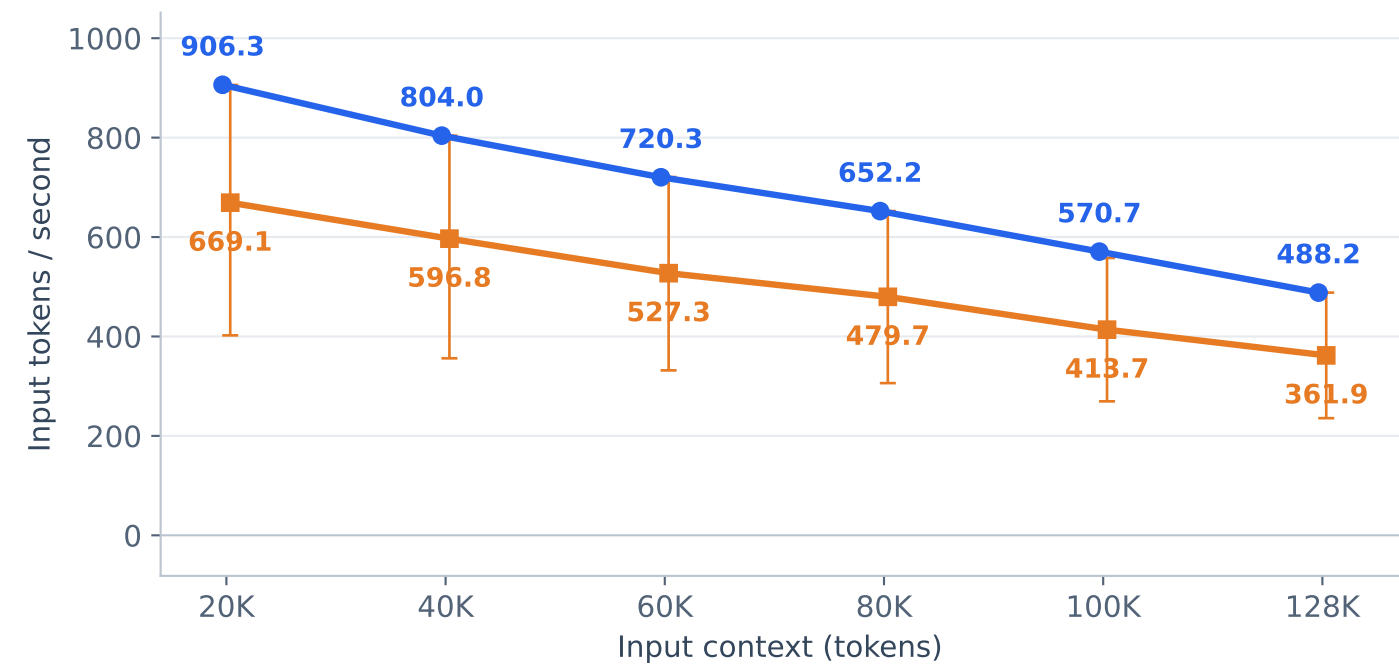
M5 Max Mac Studio · updated configuration · 1 vs 2 concurrent requests

● 1 request ■ 2 simultaneous submissions (may queue)

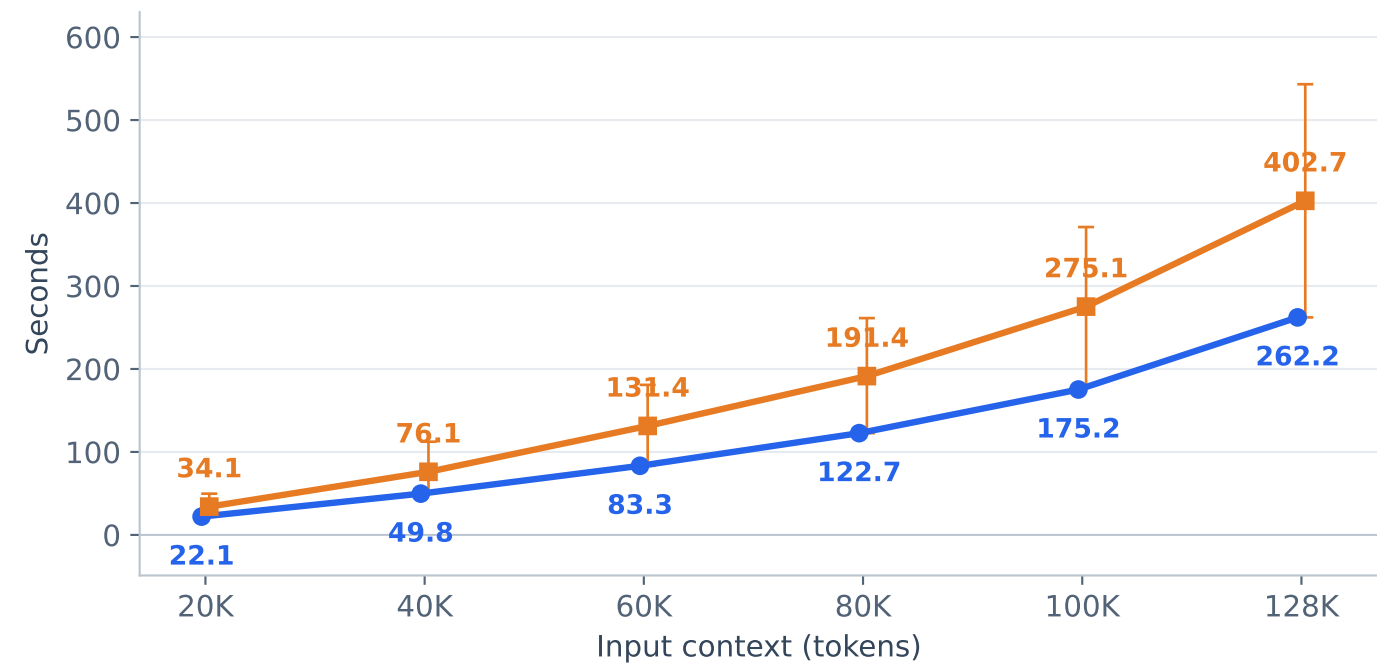
Generation speed per request



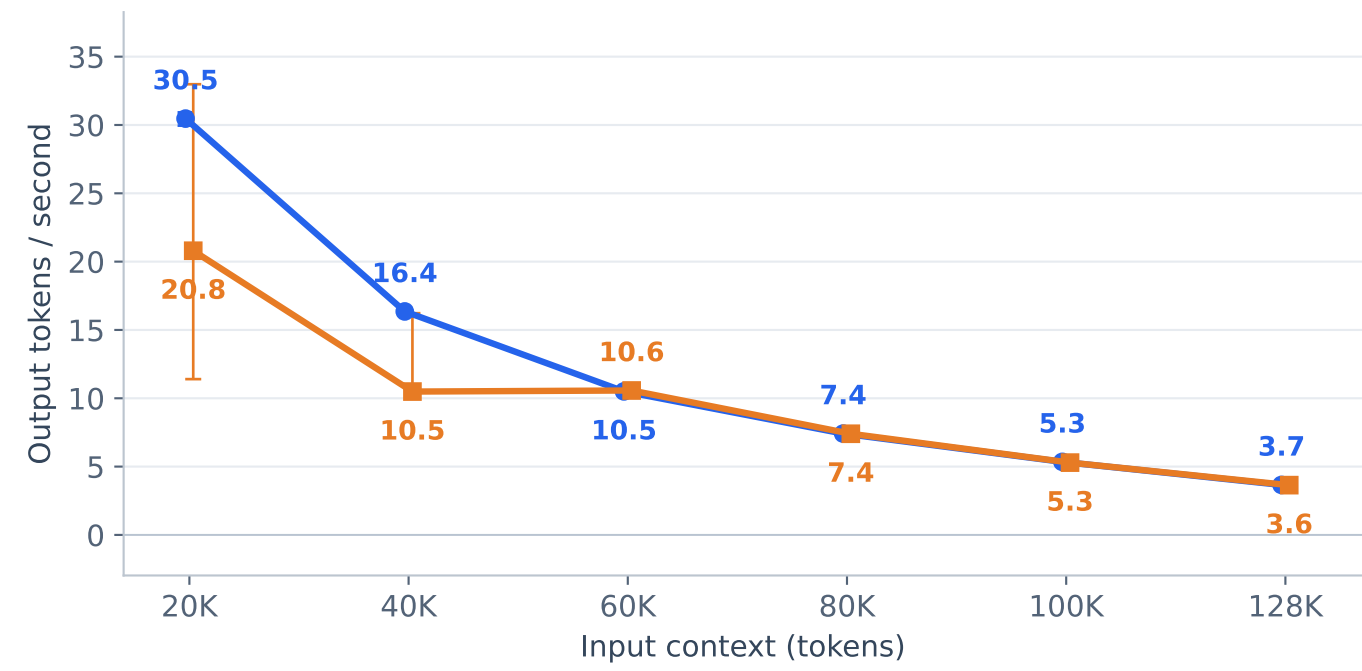
Effective prompt throughput¹



Time to first token



Combined end-to-end output throughput²



44 successful requests · 1,024 output-token limit · uncached · greedy · reasoning off · drafting on
Trials: 3 per group through 80K and 100K solo; 1 at 100K paired and 128K. Whiskers: observed min-max.
Some 20K/40K paired replies ended early (157–301 tokens); all other replies reached 1,024 tokens.
¹ Uncached input tokens ÷ (queue + prefill time), not isolated GPU prefill.
² Total output ÷ batch wall time, including queueing and prefill. Source details: updated-results.json.