


TensorFold · Swift 1.5 solo comparison


M5 Max Mac Studio · 1 request · original benchmark vs September 30

Previous Swift 1.5 configuration Updated Swift 1.5 configuration

Generation speed · higher is better



Effective prompt throughput · higher is better



Updated vs previous	20K	40K	60K	80K	100K	128K
Generation change	+10.4%	+8.5%	+8.7%	-5.9%	+6.4%	+4.8%
Prompt throughput change	+10.4%	+13.6%	+16.2%	+20.3%	+27.8%	+34.8%

Original: 3 solo trials at 20K/40K/60K; 1 at 80K/100K/128K. Current: 1 solo trial at every size.
Points: medians; whiskers: observed min-max, not confidence intervals. No repeated-trial range for September 30.
Blue: original published Swift chart, including its 128K higher-memory rerun. Pink: current 0.6.0 / M5 run.
Historical comparison, not a controlled A/B. New client uses local HTTP; earlier client used the public endpoint.
Effective prompt throughput = uncached input tokens ÷ (queue + prefill time), not isolated GPU prefill.