


TensorFold · Swift 1.5 4-bit


M5 Max Mac Studio · September 30, 2026 · 1 vs 2 concurrent requests

● 1 request ■ 2 simultaneous submissions (may queue)


Generation speed per request




Effective prompt throughput¹



Time to first token



Combined end-to-end output throughput²



18 successful requests · 1,024 output-token limit · uncached · greedy · reasoning off · drafting on

One trial per context/concurrency pair. Whiskers show the two users' range, not repeated-trial variation.

All replies reached the 1,024-token output limit; zero cached prompt tokens.

¹ Uncached input tokens ÷ (queue + prefill time), not isolated GPU prefill.

² Total output ÷ batch wall time, including queueing and prefill. Source details: updated-results.json.