


TensorFold · Swift 1.5 solo comparison


M5 Max Mac Studio · 1 request · original benchmark vs October 4

Previous Swift 1.5 configuration Updated Swift 1.5 configuration

Generation speed · higher is better



Effective prompt throughput · higher is better



Updated vs previous

Generation change

+1.6%

+14.3%

+6.5%

-6.4%

-0.6%

-0.4%

Prompt throughput change

+10.1%

+13.1%

+15.6%

+19.6%

+27.0%

+33.7%

Original: 3 solo trials at 20K/40K/60K; 1 at 80K/100K/128K. Current: 1 solo trial at every size.

Points: medians; whiskers: observed min-max, not confidence intervals. No repeated-trial range for October 4.

Blue: original published Swift chart, including its 128K higher-memory rerun. Pink: current 0.6.5 / MLX 0.32.3 / M5 run.

Historical comparison, not a controlled A/B. New client uses local HTTP; earlier client used the public endpoint.

Effective prompt throughput = uncached input tokens ÷ (queue + prefill time), not isolated GPU prefill.