TensorFold · Swift 1.5 solo comparison

M5 Max Mac Studio · 1 request · original benchmark vs October 19

Original Swift 1.5 · Current Swift 1.5 · native 1.1.1

Generation speed · higher is better

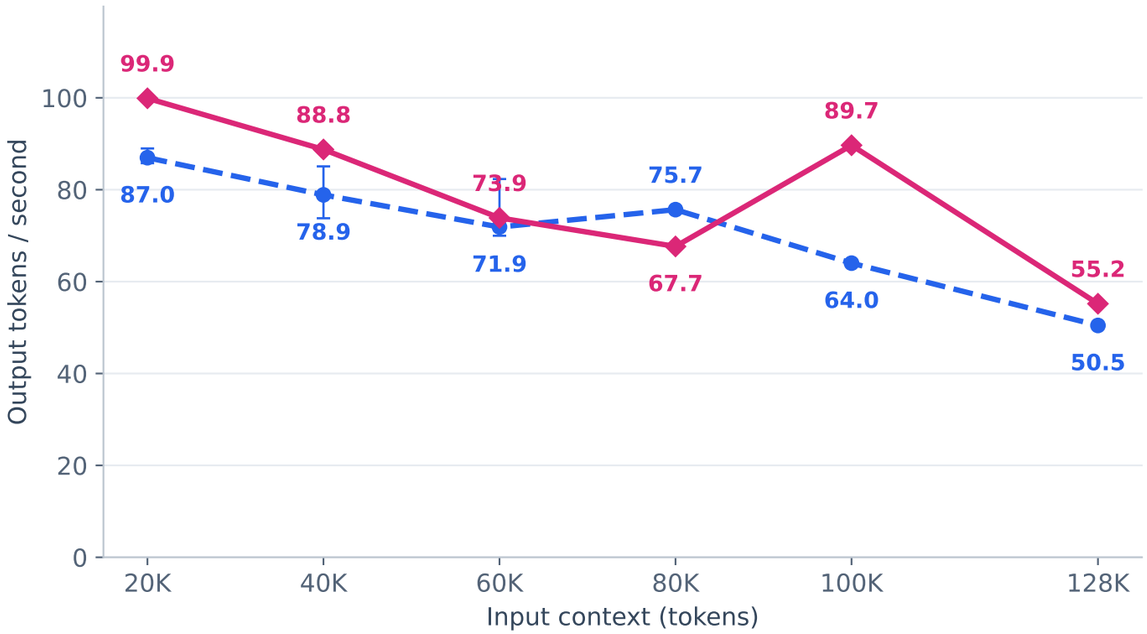

Effective prompt throughput · higher is better

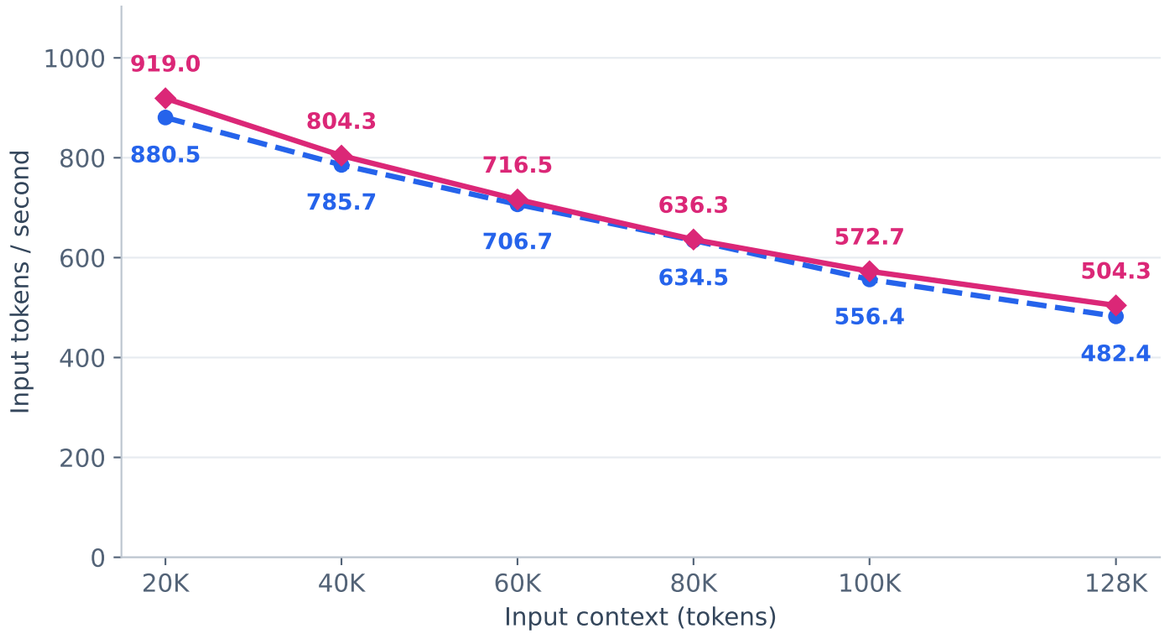

Current vs original	20K	40K	60K	80K	100K	128K
Generation change	**+14.9%**	**+12.5%**	**+2.8%**	**-10.6%**	**+40.1%**	**+9.4%**
Prompt throughput change	**+4.4%**	**+2.4%**	**+1.4%**	**+0.3%**	**+2.9%**	**+4.5%**

Original: 3 solo trials at 20K/40K/60K; 1 at 80K/100K/128K. Current: 1 solo trial at every size.

Points: medians; whiskers: observed min–max, not confidence intervals. No repeated-trial range for October 19.

Blue: original published Swift chart, including its 128K higher-memory rerun. Pink: current native 1.1.1 / Metal / M5 run.

Historical comparison, not a controlled A/B. Native vs Python engines; local vs public HTTP. Queue time restored from native receipts.

Effective prompt throughput = uncached input tokens ÷ (queue + prefill time), not isolated GPU prefill.